

TCP Reliability and Congestion Control

GW CSCI 3907/6907 Adv Networking and Distributed Systems
Prof. Timothy Wood

TCP Properties

Reliability: checksums

- Uses a 16 bit hash calculated over header/data as checksum
- Receiver can calculate checksum and verify it matches what is stored in the packet
- Is a checksum perfect?

What to do if checksum doesn't match?

TCP Segment Header Format								
Bit #	0	7	8	15	16	23	24	31
0	Source Port				Destination Port			
32	Sequence Number							
64	Acknowledgment Number							
96	Data Offset	Res	Flags		Window Size			
128	Header and Data Checksum				Urgent Pointer			
160...	Options							

TCP Properties

Reliability: based on sequence numbers and ACKs

- Client/server start connection with a random sequence number
- On every send, add the total amount of data transmitted
- On receive, reply with ACK specifying last seq number received


What to do...

- If no ACK received?
- If receiver

TCP Segment Header Format								
Bit #	0	7	8	15	16	23	24	31
0	Source Port				Destination Port			
32	Sequence Number							
64	Acknowledgment Number							
96	Data Offset	Res	Flags			Window Size		
128	Header and Data Checksum				Urgent Pointer			
160...	Options							

What happens?


Slide adapted from.
Computer Networking: A Top Down Approach
March 2012
Copyright J.F Kurose and K.W. Ross,
All Rights Reserved




ACK lost

What happens?

Slide adapted from.
Computer Networking: A Top Down Approach
March 2012
Copyright J.F Kurose and K.W. Ross,
All Rights Reserved



ACK lost



timeout and resend!
(same if original packet lost)

Wait for ACKs?


Should the **sender** wait for an ACK after each packet before sending another one?

Benefits / Drawbacks?

Wait for ACKs?

Should the **sender** wait for an ACK after each packet before sending another one?

Benefits / Drawbacks?



(a) a stop-and-wait protocol in operation

(b) a pipelined protocol in operation

Wait for ACKs?

Should the **sender** wait for an ACK after each packet before sending another one?

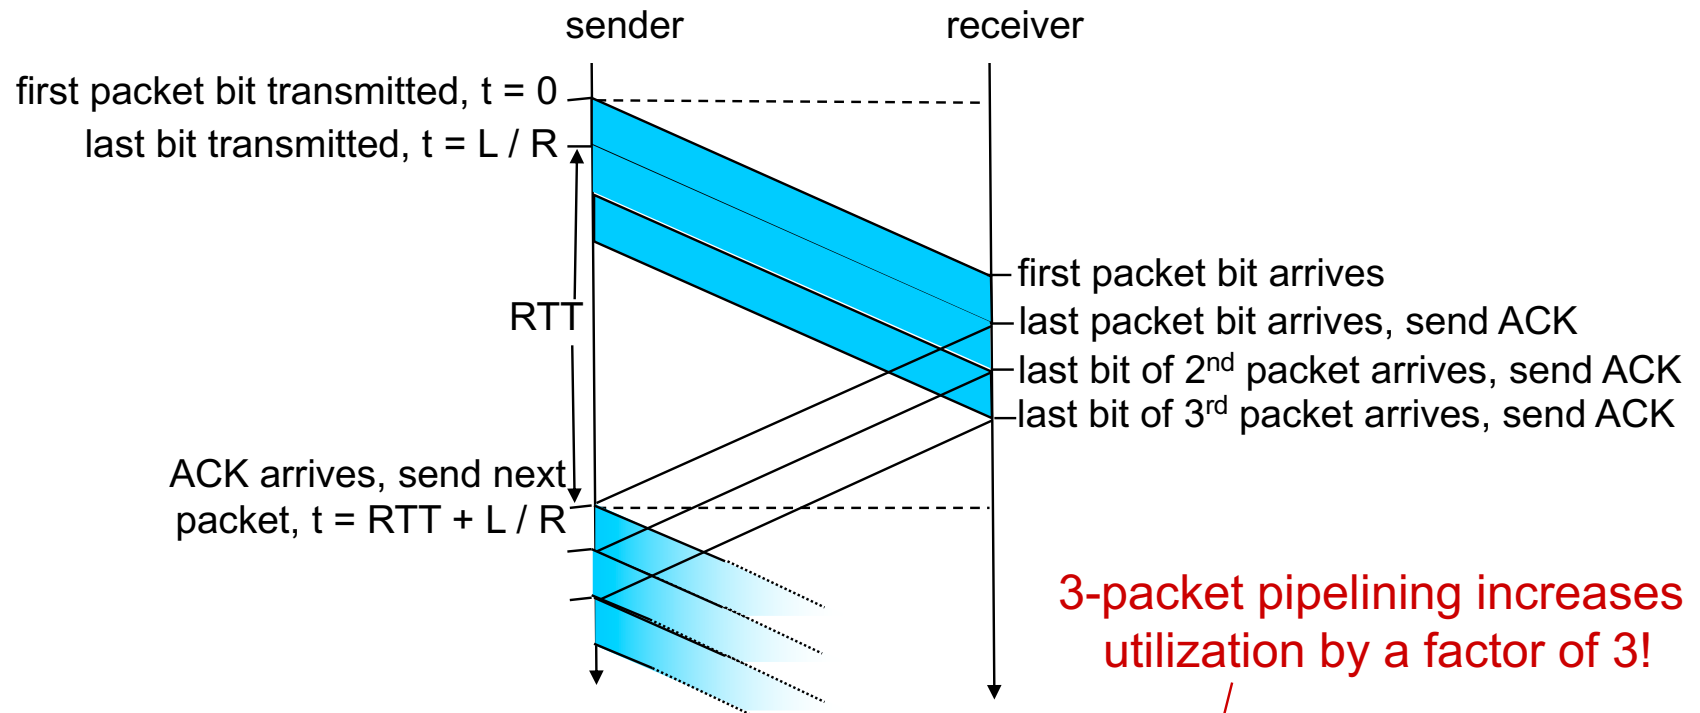
Benefits / Drawbacks?

Do the math!

Pipelining Sends

Waiting for each ACK makes very poor use of our available bandwidth!

- Better to send a “window” of packets as a pipeline



3-packet pipelining increases utilization by a factor of 3!

$$U_{\text{sender}} = \frac{3L/R}{RTT + L/R} = \frac{.0024}{30.008} = 0.00081$$


Slide adapted from.

Computer Networking: A Top Down Approach
March 2012

Copyright J.F Kurose and K.W. Ross,
All Rights Reserved

What happens?


Slide adapted from.
Computer Networking: A Top Down Approach
March 2012
Copyright J.F Kurose and K.W. Ross,
All Rights Reserved



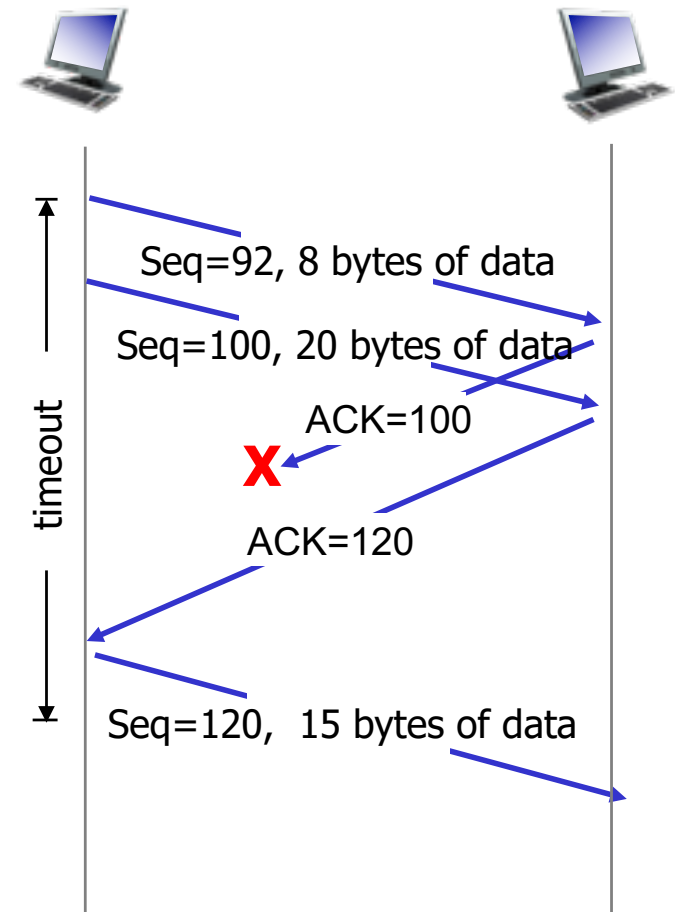
1st ACK lost

What happens?

Slide adapted from.
Computer Networking: A Top Down Approach
March 2012
Copyright J.F Kurose and K.W. Ross,
All Rights Reserved



1st ACK lost

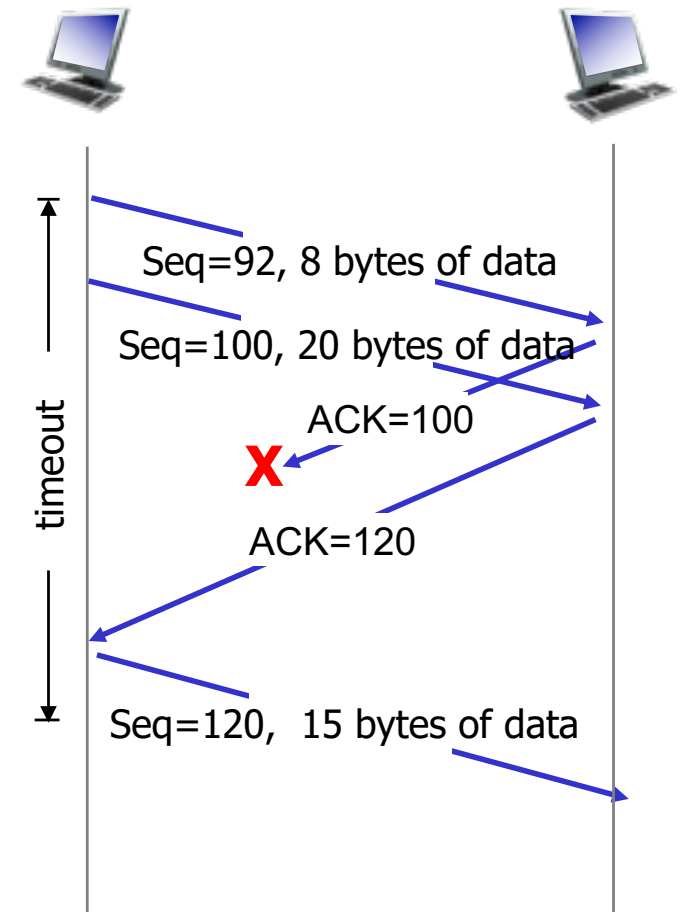


2nd ACK is
cumulative!

Cumulative ACKs

ACK 120 means ALL bytes up to that point are received

Why use cumulative instead of individual ACKs?



2nd ACK is cumulative!

Wait for ACKs?

Should the **receiver** immediately send an ACK?

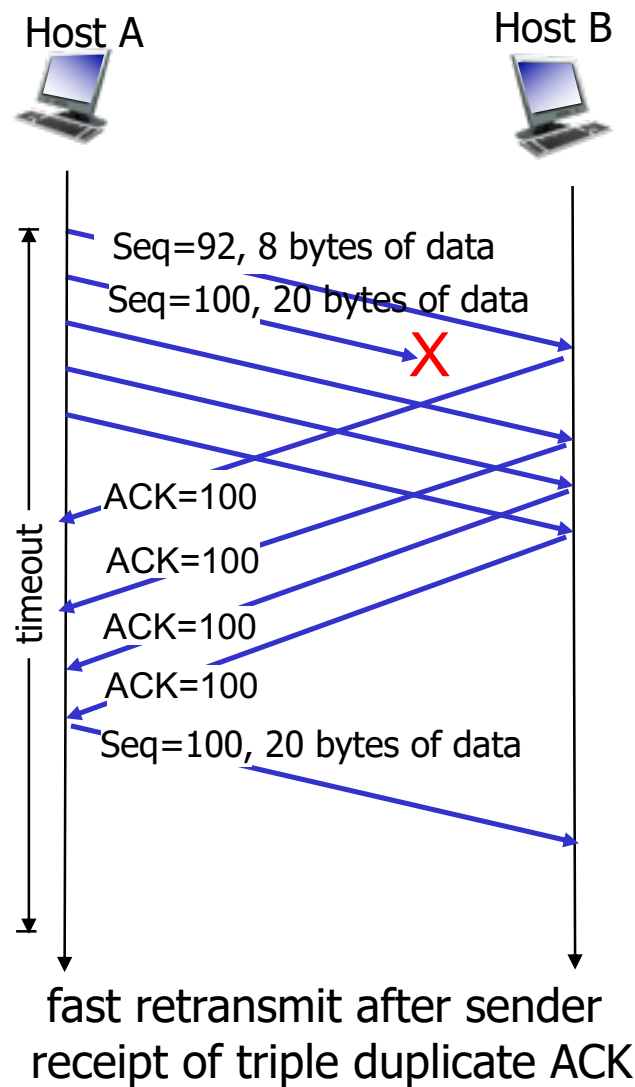
Benefits / Drawbacks?

TCP Reliability

Slide adapted from.
Computer Networking: A Top Down Approach
March 2012
Copyright J.F Kurose and K.W. Ross,
All Rights Reserved

<i>event at receiver</i>	<i>TCP receiver action</i>
arrival of in-order segment with expected seq #. All data up to expected seq # already ACKed	delayed ACK. Wait up to 500ms for next segment. If no next segment, send ACK
arrival of in-order segment with expected seq #. One other segment has ACK pending	immediately send single cumulative ACK, ACKing both in-order segments
arrival of out-of-order segment higher-than-expect seq. # . Gap detected	immediately send <i>duplicate ACK</i> , indicating seq. # of next expected byte
arrival of segment that partially or completely fills gap	immediate send ACK, provided that segment starts at lower end of gap

TCP Fast Retransmit



Slide adapted from.
Computer Networking: A Top Down Approach
March 2012
Copyright J.F Kurose and K.W. Ross,
All Rights Reserved

How many packets to send?

Using a larger **window** leads to better link utilization

- *So why not just use a window of 1,000,000?*

Benefit of large window?

Drawback of large window?

How many packets to send?

Using a larger **window** leads to better link utilization

- *So why not just use a window of 1,000,000?*

Benefit of large window?

- Can send more data before waiting for ACK

Drawback of large window?

- With cumulative ACK, it can be hard for sender to know exactly what it needs to resend
- Sending at too high a rate may cause higher packet loss!


This leads to why TCP does Congestion Control!

- We'll just cover the basics...

Windows

Window size controls # of packets in flight


**Sender's
view of
Sequence
Numbers**



Windows

Window size controls # of packets in flight

**Sender's
view of
Sequence
Numbers**



**Receiver's
view of
Sequence
Numbers**



Remember
that TCP is
bidirectional,
so this all
happens
twice!

Congestion Control Basics

How should we adjust window size?

- Let's assume client is sending a large file to server

Congestion Control Basics

How should we adjust window size?

- Let's assume client is sending a large file to server

Startup:

- use a small window since you don't know anything about receiver

No drops for a while:

```
window_size++; // Send faster!
```

Drop detected:

```
window_size = window_size/2; // whoa! slow down!
```

Additive Increase, Multiplicative Decrease (AIMD)